

Encoder baselines — frozen backbone → pooling → linear head

Whisper × mean pool



Whisper × attention pool



WavLM × mean pool



WavLM × attention pool



Fused encoder (Whisper + WavLM)

